

Loan Default Prediction

Team 17

Xin Feng Kaimi Huang

Yuanchen Cui Boqian Wang



CONTENT

01

Project Introduction

02

**Attempts to Improve
Model Performance**

03

**Gaussian Mixture
Clusters**

04

**Results Analysis and
Conclusion**

Project Introduction

Loan default rate prediction is **a critical aspect of financial risk management** for banks, credit unions, and other lending institutions.

01




The ability to accurately predict loan default rate can significantly **impact the financial health of these institutions.**

02

It enables lenders to **mitigate risks, set appropriate interest rates, allocate reserves for potential losses, and make informed decisions** about loan approvals.

03

It also helps in **tailoring loan products** to suit different risk profiles, **enhancing customer satisfaction**, and **fostering financial stability** in the broader economy.



Data Collection and Preprocessing

Data Source:

Lending Club Loan Dataset

(<https://www.scaler.com/topics/data-science/loan-default-prediction/>)

20000 records, 15 columns

Target:

bad_loan

Features:

id、grade、annual_income、
short_employee、emp_length_num、
home_ownership、Debt-To-Income Ratio、
purpose、term、last_delinq_none、
last_major_derog_none、revol_util、
total_rec_late_fee、od_ratio、bad loan

Missing Values:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 20000 entries, 0 to 19999
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	id	20000 non-null	int64
1	grade	20000 non-null	object
2	annual_income	20000 non-null	int64
3	short_employee	20000 non-null	int64
4	emp_length_num	20000 non-null	int64
5	home_ownership	18509 non-null	object
6	Debt-To-Income Ratio	19846 non-null	float64
7	purpose	20000 non-null	object
8	term	20000 non-null	object
9	last_delinq_none	20000 non-null	int64
10	last_major_derog_none	574 non-null	float64
11	revol_util	20000 non-null	float64
12	total_rec_late_fee	20000 non-null	float64
13	od_ratio	20000 non-null	float64
14	<u>bad_loan</u>	20000 non-null	int64



```
dtypes: float64(5), int64(6), object(4)
```

```
memory usage: 2.3+ MB
```

Outlier

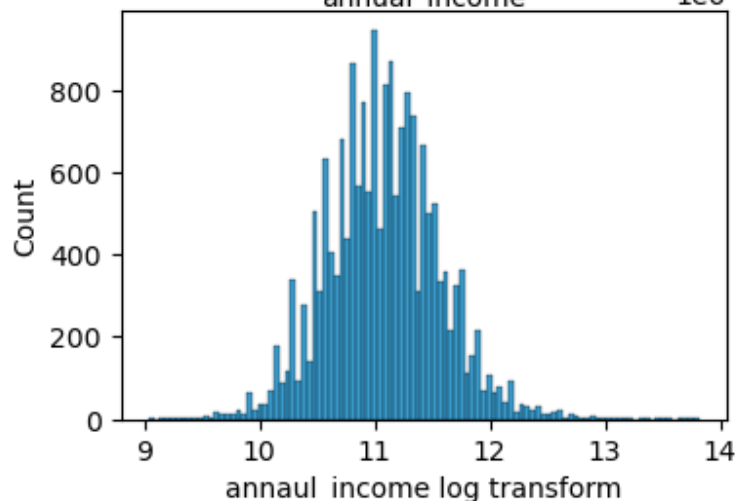

```
df.describe()
```

	id	annual_income	short_employee	emp_length_num	Debt-To-Income Ratio	revol_util	total_rec_late_fee	od_ratio	bad_loan
count	1.837100e+04	18371.000000	18371.000000	18371.000000	18371.000000	18371.000000	18371.000000	18371.000000	18371.000000
mean	7.594628e+06	73421.273257	0.112297	6.827609	16.590894	<u>56.001801</u>	<u>0.293404</u>	0.504941	0.200479
std	1.609952e+06	45612.958798	0.315740	3.769322	7.582902	43.411698	3.140913	0.287800	0.400370
min	5.860400e+05	8412.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000077	0.000000
25%	6.206280e+06	47000.000000	0.000000	3.000000	10.850000	38.750000	0.000000	0.257495	0.000000
50%	7.379923e+06	65000.000000	0.000000	7.000000	16.220000	57.100000	0.000000	0.507883	0.000000
75%	8.776061e+06	88000.000000	0.000000	11.000000	22.060000	74.000000	0.000000	0.753875	0.000000
max	1.145464e+07	100000.000000	1.000000	11.000000	34.990000	<u>5010.000000</u>	<u>96.466600</u>	0.999894	1.000000



Feature Transformation & Engineering

1. Transform annual income to log (annual income)

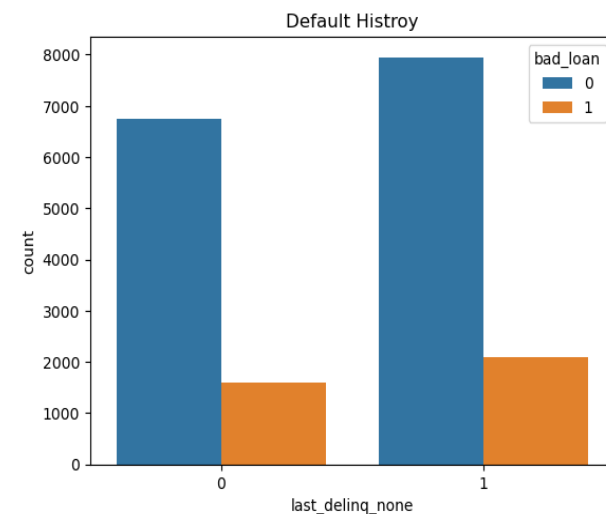


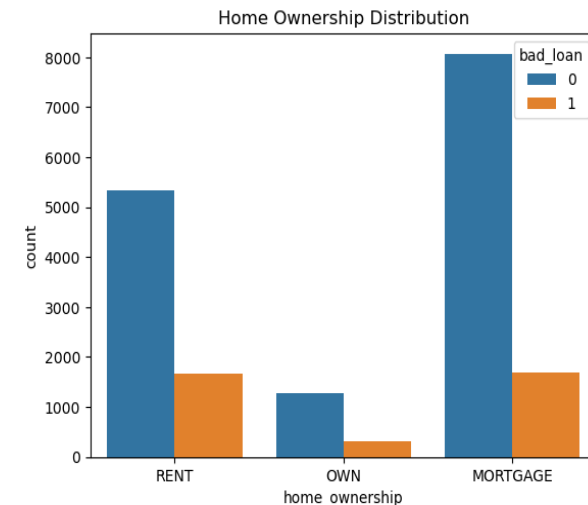




2. One-Hot Encoding for Categorical Features

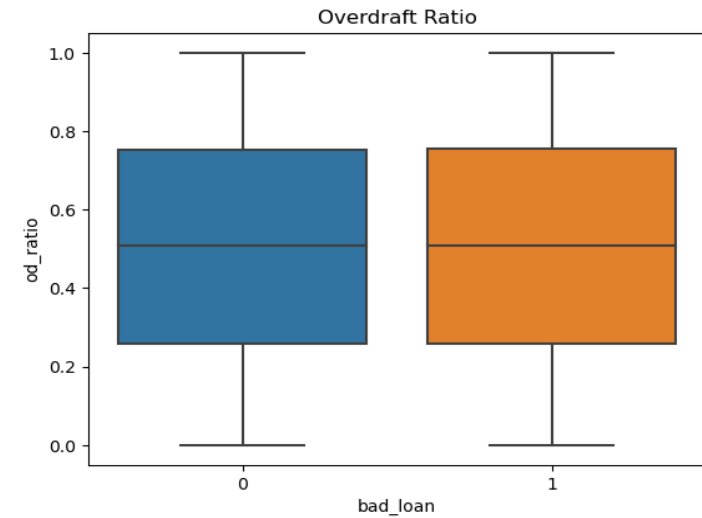
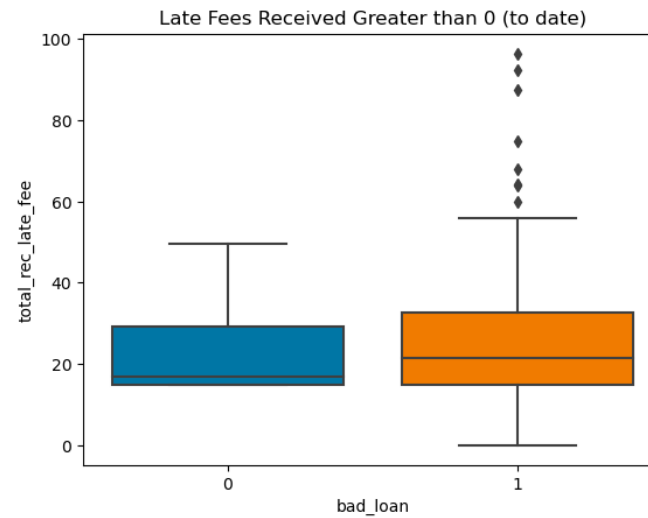
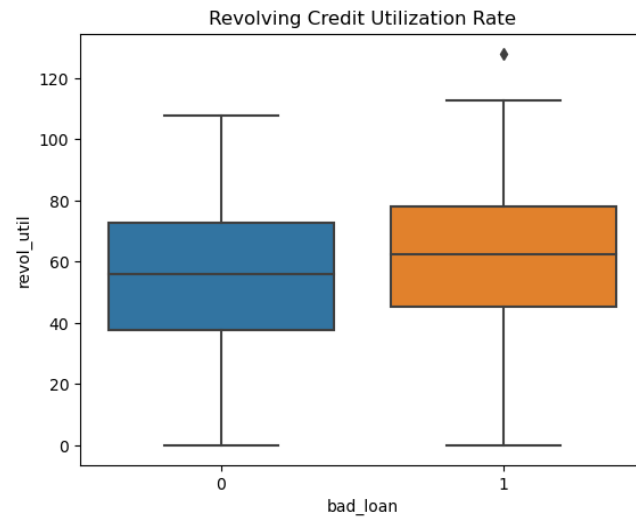
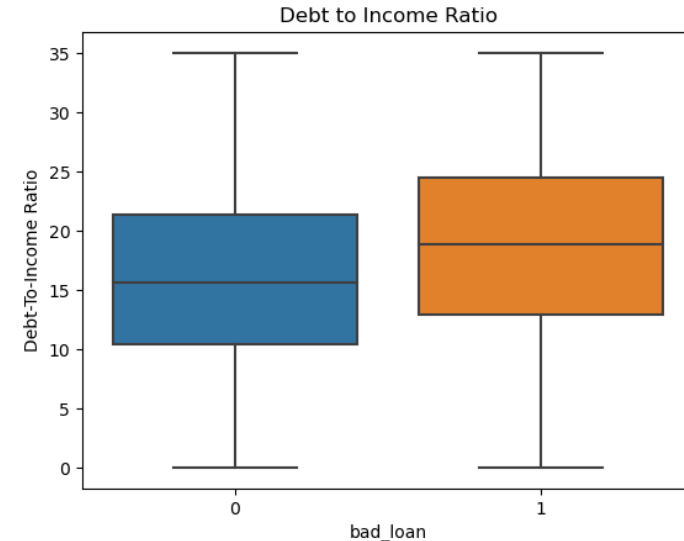


```
1 X.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 18370 entries, 0 to 18369  
Data columns (total 28 columns):  
#   Column                                Non-Null Count  Dtype  
---  ---                                -  
0   short_employee                        18370 non-null  int64  
1   emp_length_num                       18370 non-null  int64  
2   Debt-To-Income Ratio                 18370 non-null  float64  
3   revol_util                           18370 non-null  float64  
4   total_rec_late_fee                   18370 non-null  float64  
5   od_ratio                             18370 non-null  float64  
6   annual_income_log                    18370 non-null  float64  
7   grade_B                              18370 non-null  uint8  
8   grade_C                              18370 non-null  uint8  
9   grade_D                              18370 non-null  uint8  
10  grade_E                              18370 non-null  uint8  
11  grade_F                              18370 non-null  uint8  
12  grade_G                              18370 non-null  uint8  
13  home_ownership_OWNI                  18370 non-null  uint8  
14  home_ownership_RENT                  18370 non-null  uint8  
15  purpose_credit_card                  18370 non-null  uint8  
16  purpose_debt_consolidation           18370 non-null  uint8  
17  purpose_home_improvement             18370 non-null  uint8  
18  purpose_house                        18370 non-null  uint8  
19  purpose_major_purchase               18370 non-null  uint8  
20  purpose_medical                      18370 non-null  uint8  
21  purpose_moving                       18370 non-null  uint8  
22  purpose_other                        18370 non-null  uint8  
23  purpose_small_business               18370 non-null  uint8  
24  purpose_vacation                     18370 non-null  uint8  
25  purpose_wedding                      18370 non-null  uint8  
26  term_60 months                       18370 non-null  uint8  
27  last_delinq_none_1                  18370 non-null  uint8  
dtypes: float64(5), int64(2), uint8(21)
```

Distributions of Target and Categorical Features



Distributions of Continuous Features



Training Models

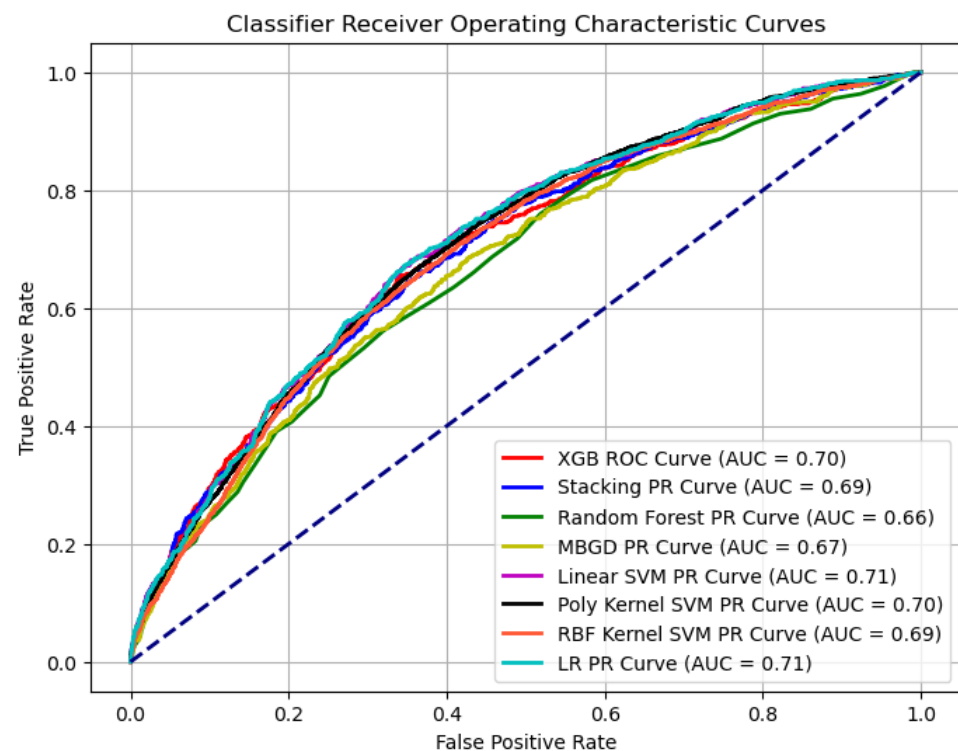
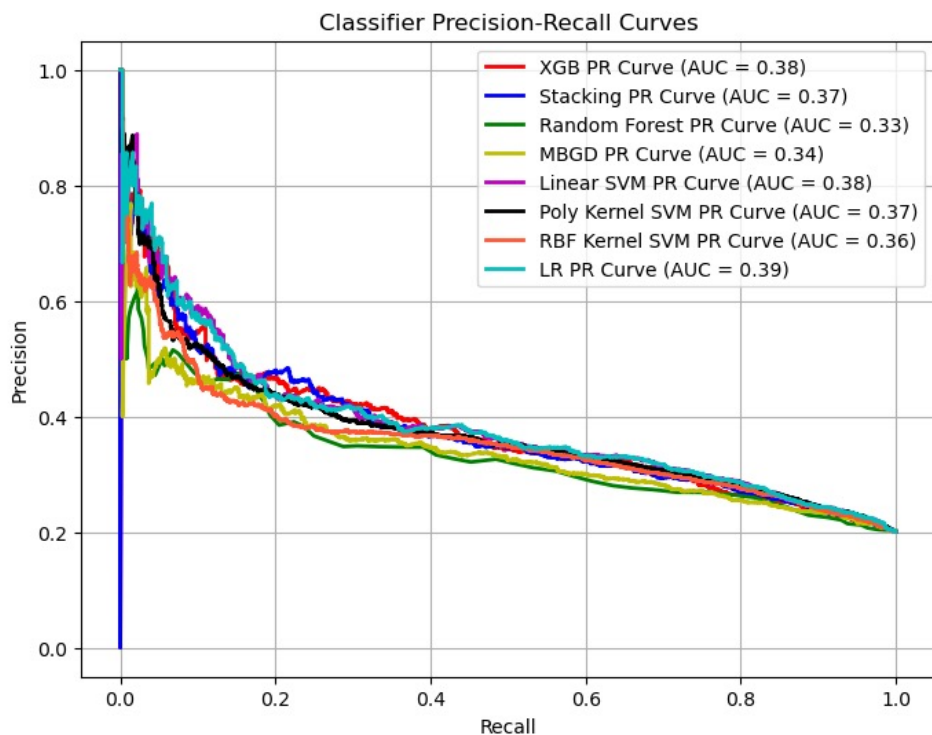
- `X = df.drop(columns=['id', 'annual_income','bad_loan','grade_A',
 'home_ownership_MORTGAGE','purpose_car','term_36 months',
 'last_delinq_none_0'])`
- `y = df['bad_loan'].astype(int)` *#1 represents bad loan. 0 good loan.*
- `StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=123)`
- `scaler = StandardScaler()`
- `X_train = scaler.fit_transform(X_train)`
- `X_test = scaler.transform(X_test)`
- `class_weight='balanced'` *#hyperparameter for all classifier instances*
- `GridSearchCV()` *#used for finding optimal hyperparameters for all classifiers*

```
1 X.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 18370 entries, 0 to 18369  
Data columns (total 28 columns):  
#   Column                                Non-Null Count  Dtype  
---  -  
0   short_employee                        18370 non-null  int64  
1   emp_length_num                       18370 non-null  int64  
2   Debt-To-Income Ratio                 18370 non-null  float64  
3   revol_util                           18370 non-null  float64  
4   total_rec_late_fee                   18370 non-null  float64  
5   od_ratio                             18370 non-null  float64  
6   annual_income_log                    18370 non-null  float64  
7   grade_B                              18370 non-null  uint8  
8   grade_C                              18370 non-null  uint8  
9   grade_D                              18370 non-null  uint8  
10  grade_E                              18370 non-null  uint8  
11  grade_F                              18370 non-null  uint8  
12  grade_G                              18370 non-null  uint8  
13  home_ownership_OWN                   18370 non-null  uint8  
14  home_ownership_RENT                  18370 non-null  uint8  
15  purpose_credit_card                  18370 non-null  uint8  
16  purpose_debt_consolidation           18370 non-null  uint8  
17  purpose_home_improvement             18370 non-null  uint8  
18  purpose_house                        18370 non-null  uint8  
19  purpose_major_purchase               18370 non-null  uint8  
20  purpose_medical                      18370 non-null  uint8  
21  purpose_moving                      18370 non-null  uint8  
22  purpose_other                        18370 non-null  uint8  
23  purpose_small_business               18370 non-null  uint8  
24  purpose_vacation                     18370 non-null  uint8  
25  purpose_wedding                      18370 non-null  uint8  
26  term_60 months                       18370 non-null  uint8  
27  last_delinq_none_1                   18370 non-null  uint8  
dtypes: float64(5), int64(2), uint8(21)
```

Model Performances

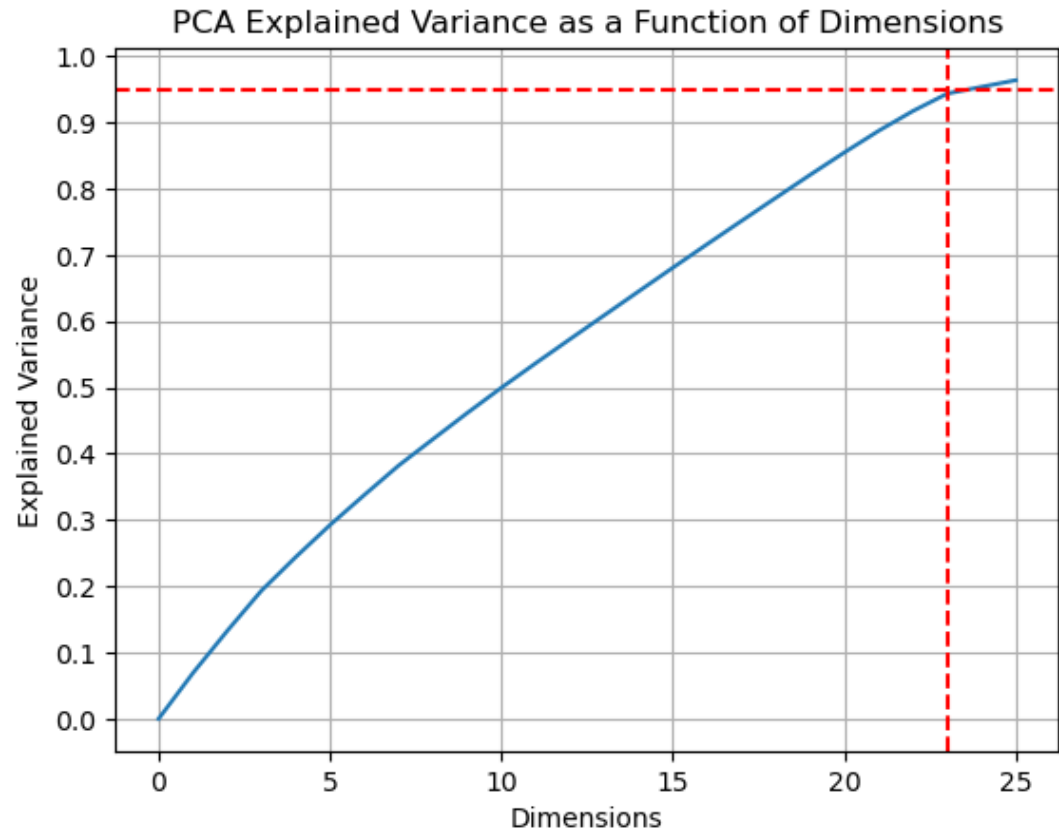
Model	PR AUC	ROC AUC	Accuracy
Logistic Regression	0.39	0.71	0.66
XGBoost	0.38	0.7	0.8
Random Forest	0.33	0.66	0.8
Stochastic Gradient Descent	0.34	0.67	0.61
Mini-batch Gradient Descent	0.34	0.67	0.61
Linear Support Vector Machine	0.38	0.71	0.66
Polynomial Kernel Support Vector Classifier	0.37	0.7	0.65
Gaussian RBF Kernel Support Vector Classifier	0.36	0.69	0.72
Stacking/Stacked Generalization	0.37	0.69	0.65



Attempts to Improve Model Performance

1. Dimensionality Reduction

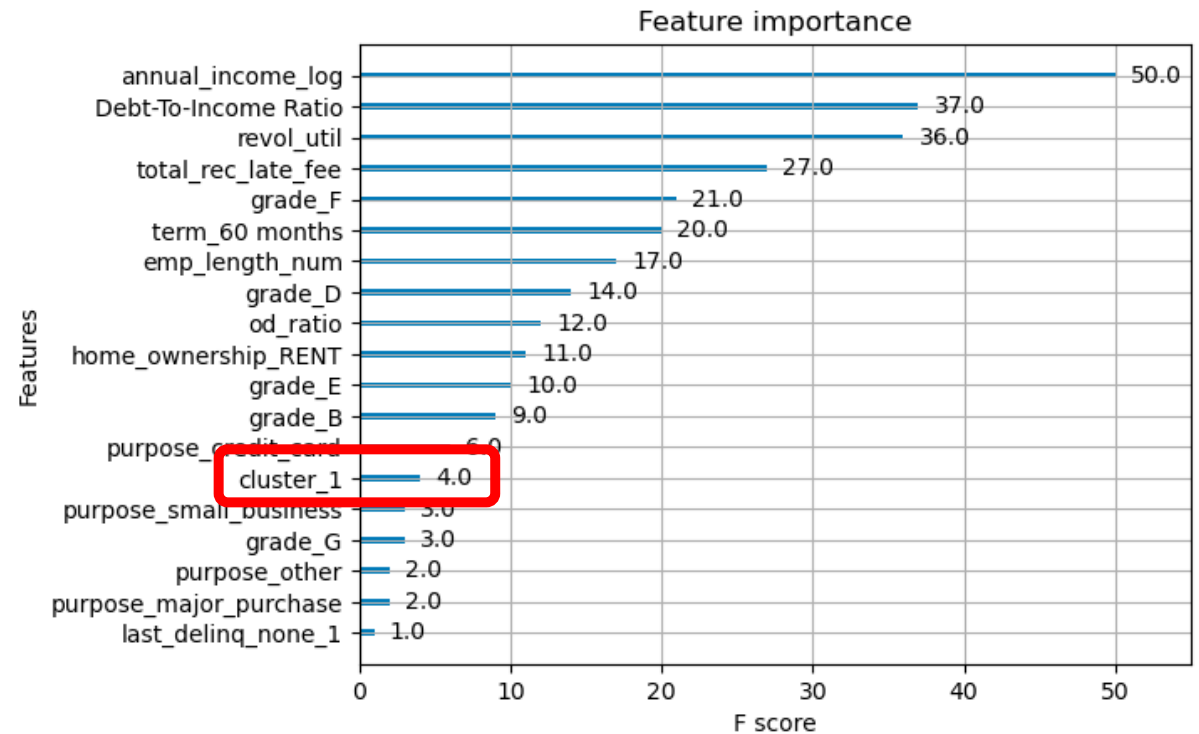
- **Idea:** reduce noise via PCA
- In order to preserve 95% of the variance, kept 23 of the 25 features
- `pca = PCA(n_components=0.95)`
- **Result:** PR AUCs did not improve but dropped slightly



Attempts to Improve Model Performance

2. Create new categorical features, “cluster_0” & “cluster_1”, using K-means

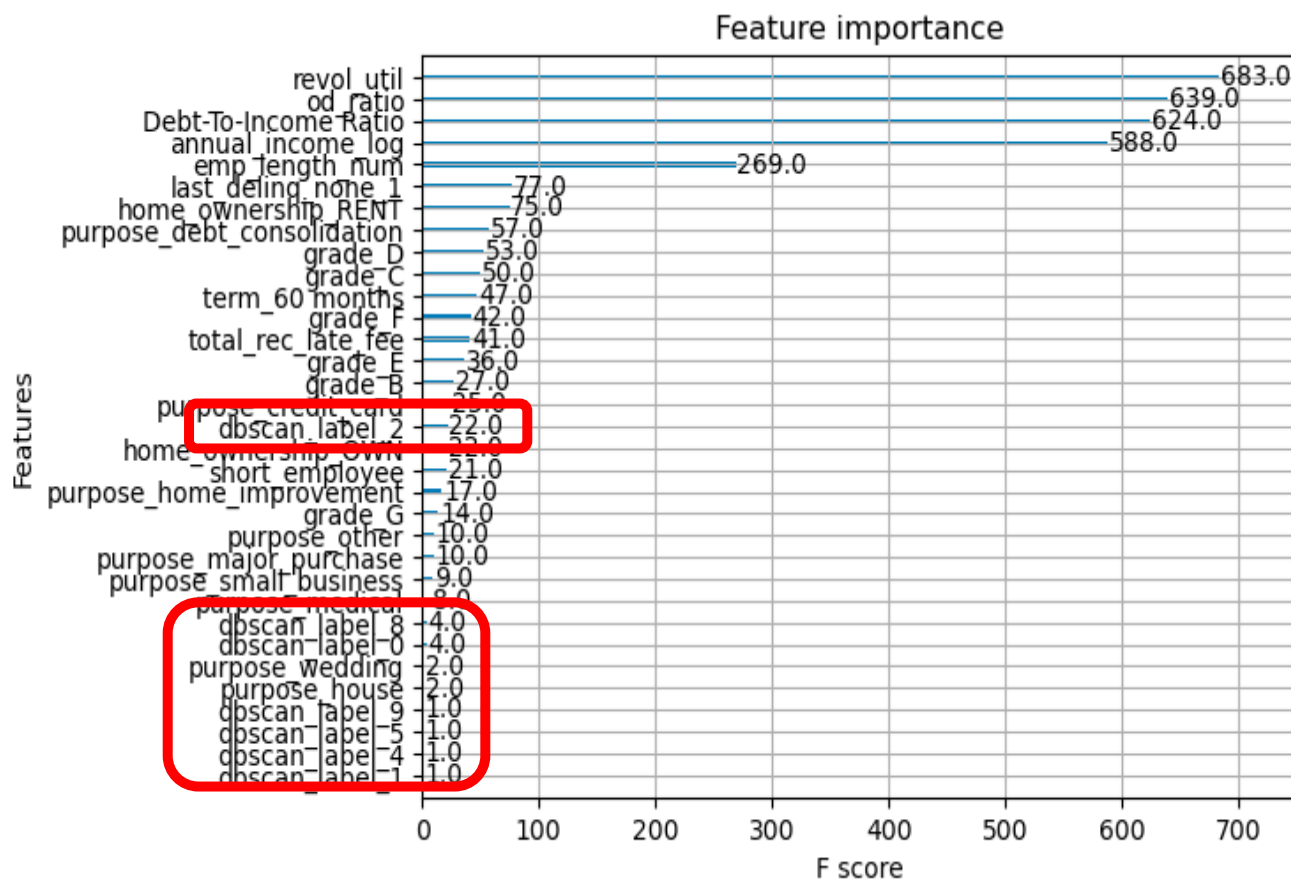
- **Idea:** grouping loans based on features might add information that help classification
- `pipeline = Pipeline([
 ('kmeans', KMeans()),
 ('log_reg', lg_clf)])`
- Grid search showed the optimal number of cluster was `{'kmeans__n_clusters': 3}`
- **Result:** PR AUCs did not improve.
New feature's importance was low.



Attempts to Improve Model Performance

3. Create new categorical features, “dbscan_label_”, using DBSCAN

- **Idea:** grouping loans based on features might add information that help classification
- `dbscan = DBSCAN(eps = 6, min_samples = 3)`
- `X_train` was clustered into 13 groups
- **Result:** PR AUC was worse.
New features' importance was low



What if the target labels were wrong?

- Separate df into two datasets by target label
- Find anomalies in each set using Gaussian Mixture
 - setting n_components=2 because of the assumption that some of the observations in the good loans are mislabeled and should have been labeled as bad loans, and vice versa

```
1 df_good_loan = df[df.bad_loan==0]
2 df_bad_loan = df[df.bad_loan==1]
```

```
1 gm_good = GaussianMixture(n_components=2, n_init=10, random_state=123)
2 gm_good.fit(df_good_loan)
```


```
GaussianMixture(n_components=2, n_init=10, random_state=123)
```

```
1 gm_good.weights_
```

```
array([0.06910874, 0.93089126])
```


```
1 gm_bad = GaussianMixture(n_components=2, n_init=20, random_state=123)
2 gm_bad.fit(df_bad_loan)
3 gm_bad.weights_
```

```
array([0.18028848, 0.81971152])
```



Gaussian Mixture Clusters


Original label = 0



Correctly labeled
as 0

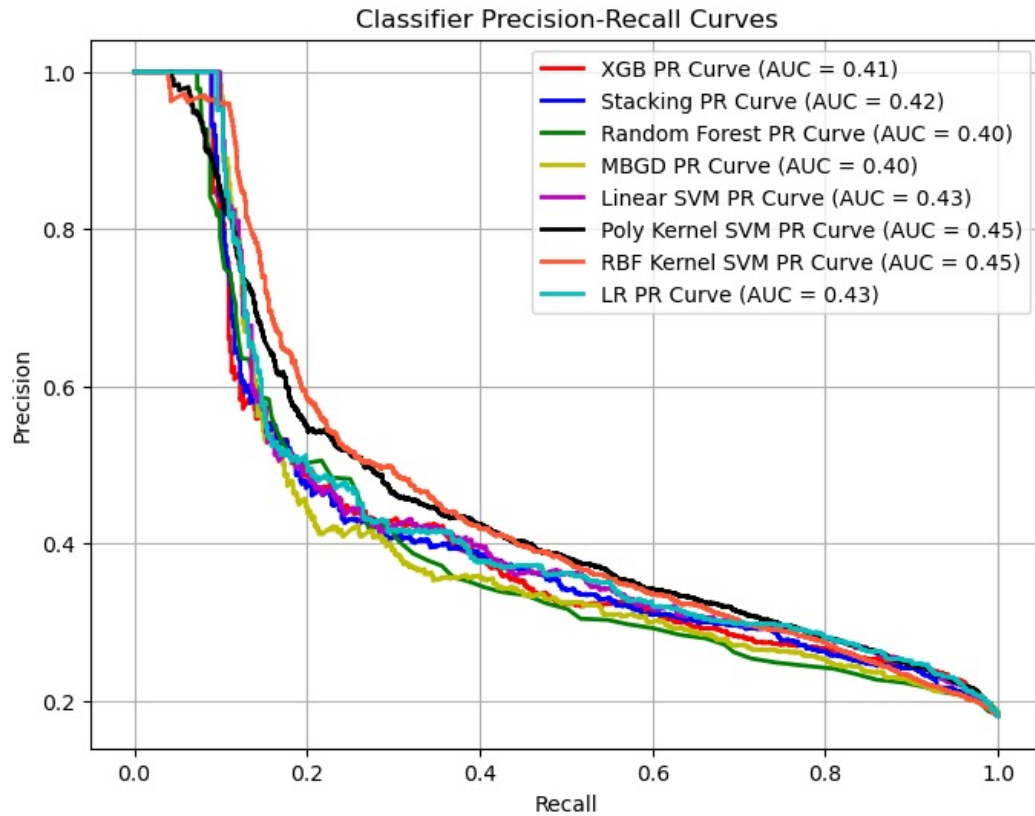
Does it mean the labels of these two
bigger clusters should be flipped?

Original label = 1

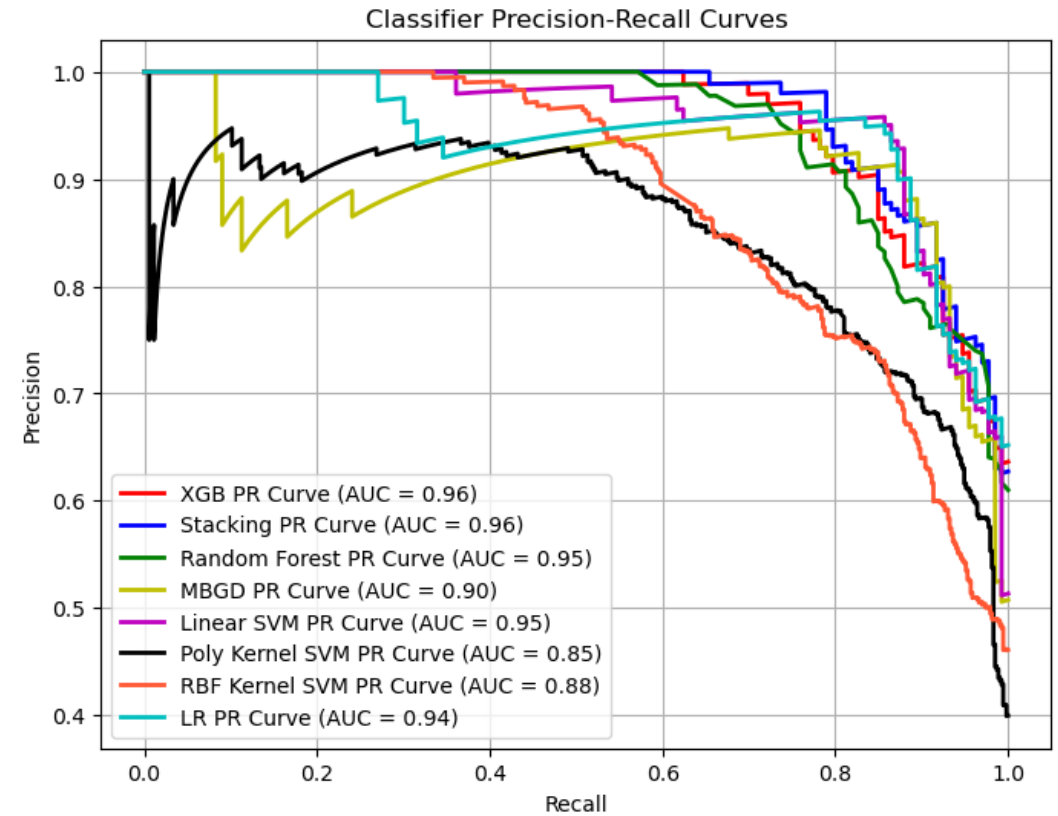


Correctly labeled
as 1

Model Performances using the bigger two sets of data

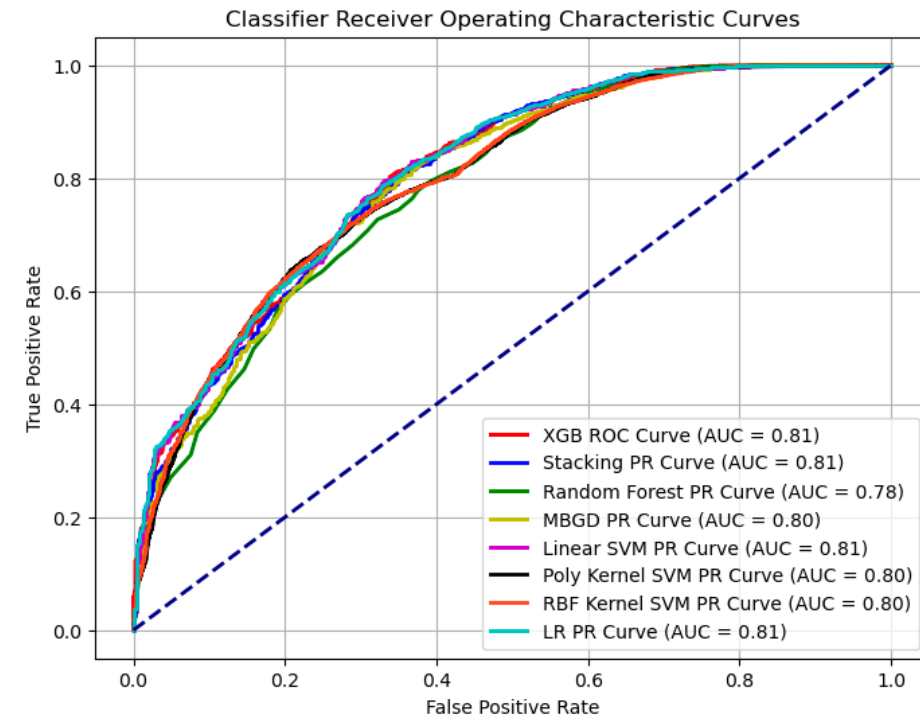



Model Performances using the smaller two sets of data

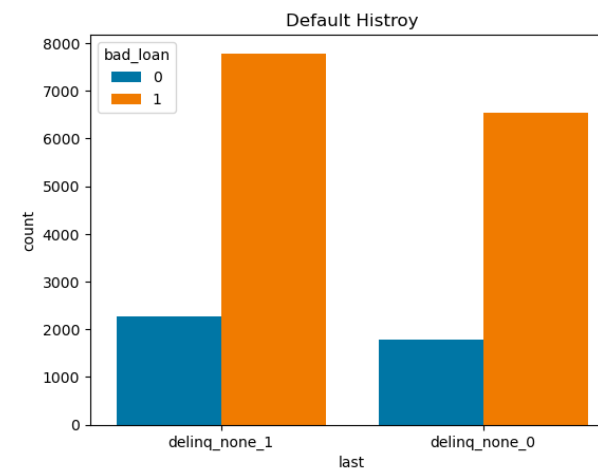







Model Performances after Flipping Labels

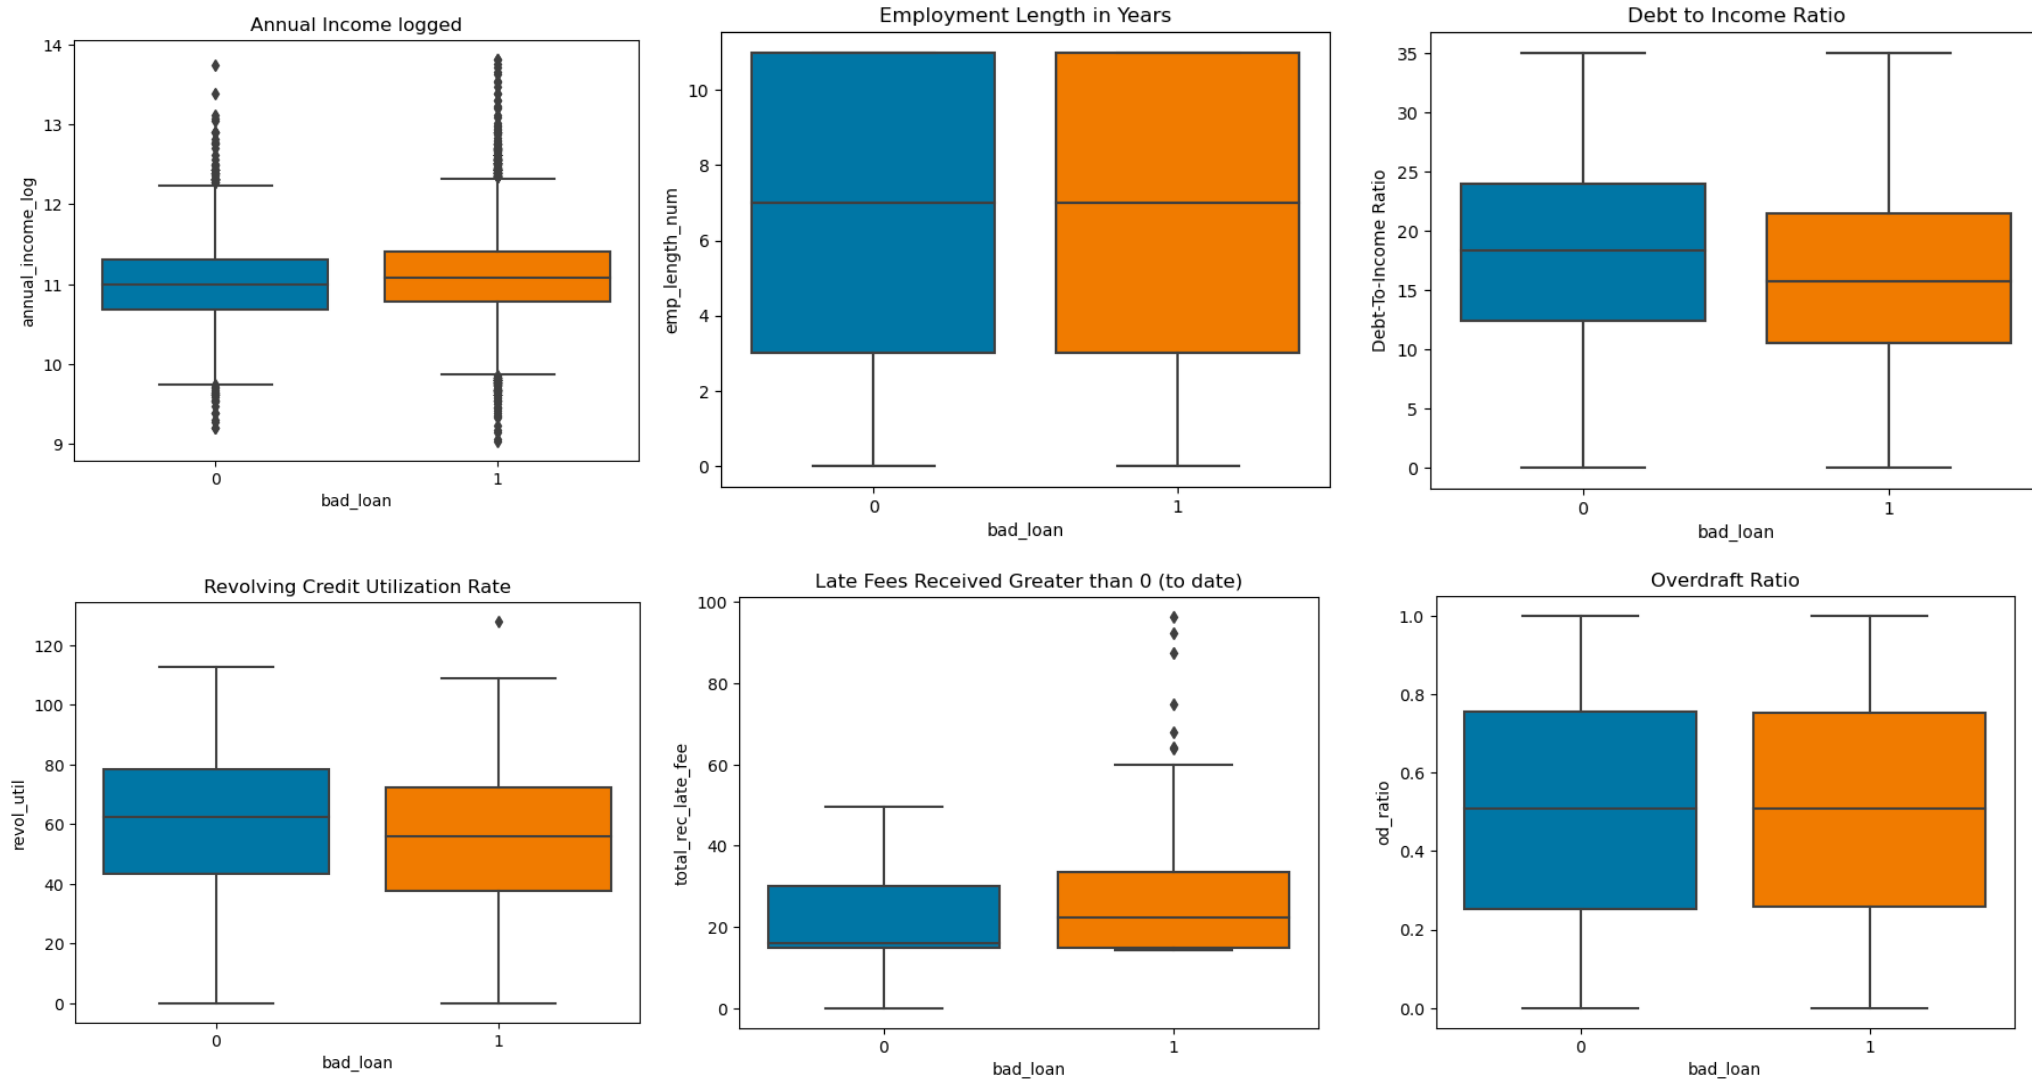
Model	PR AUC	ROC AUC	Accuracy
Logistic Regression	0.93	0.81	0.76
XGBoost	0.93	0.81	0.84
Random Forest	0.92	0.78	0.83
Stochastic Gradient Descent	0.92	0.8	0.75
Mini-batch Gradient Descent	0.92	0.8	0.75
Linear Support Vector Machine	0.93	0.81	0.77
Polynomial Kernel Support Vector Classifier	0.92	0.8	0.76
Gaussian RBF Kernel Support Vector Classifier	0.92	0.8	0.76
Stacking/Stacked Generalization	0.93	0.81	0.76

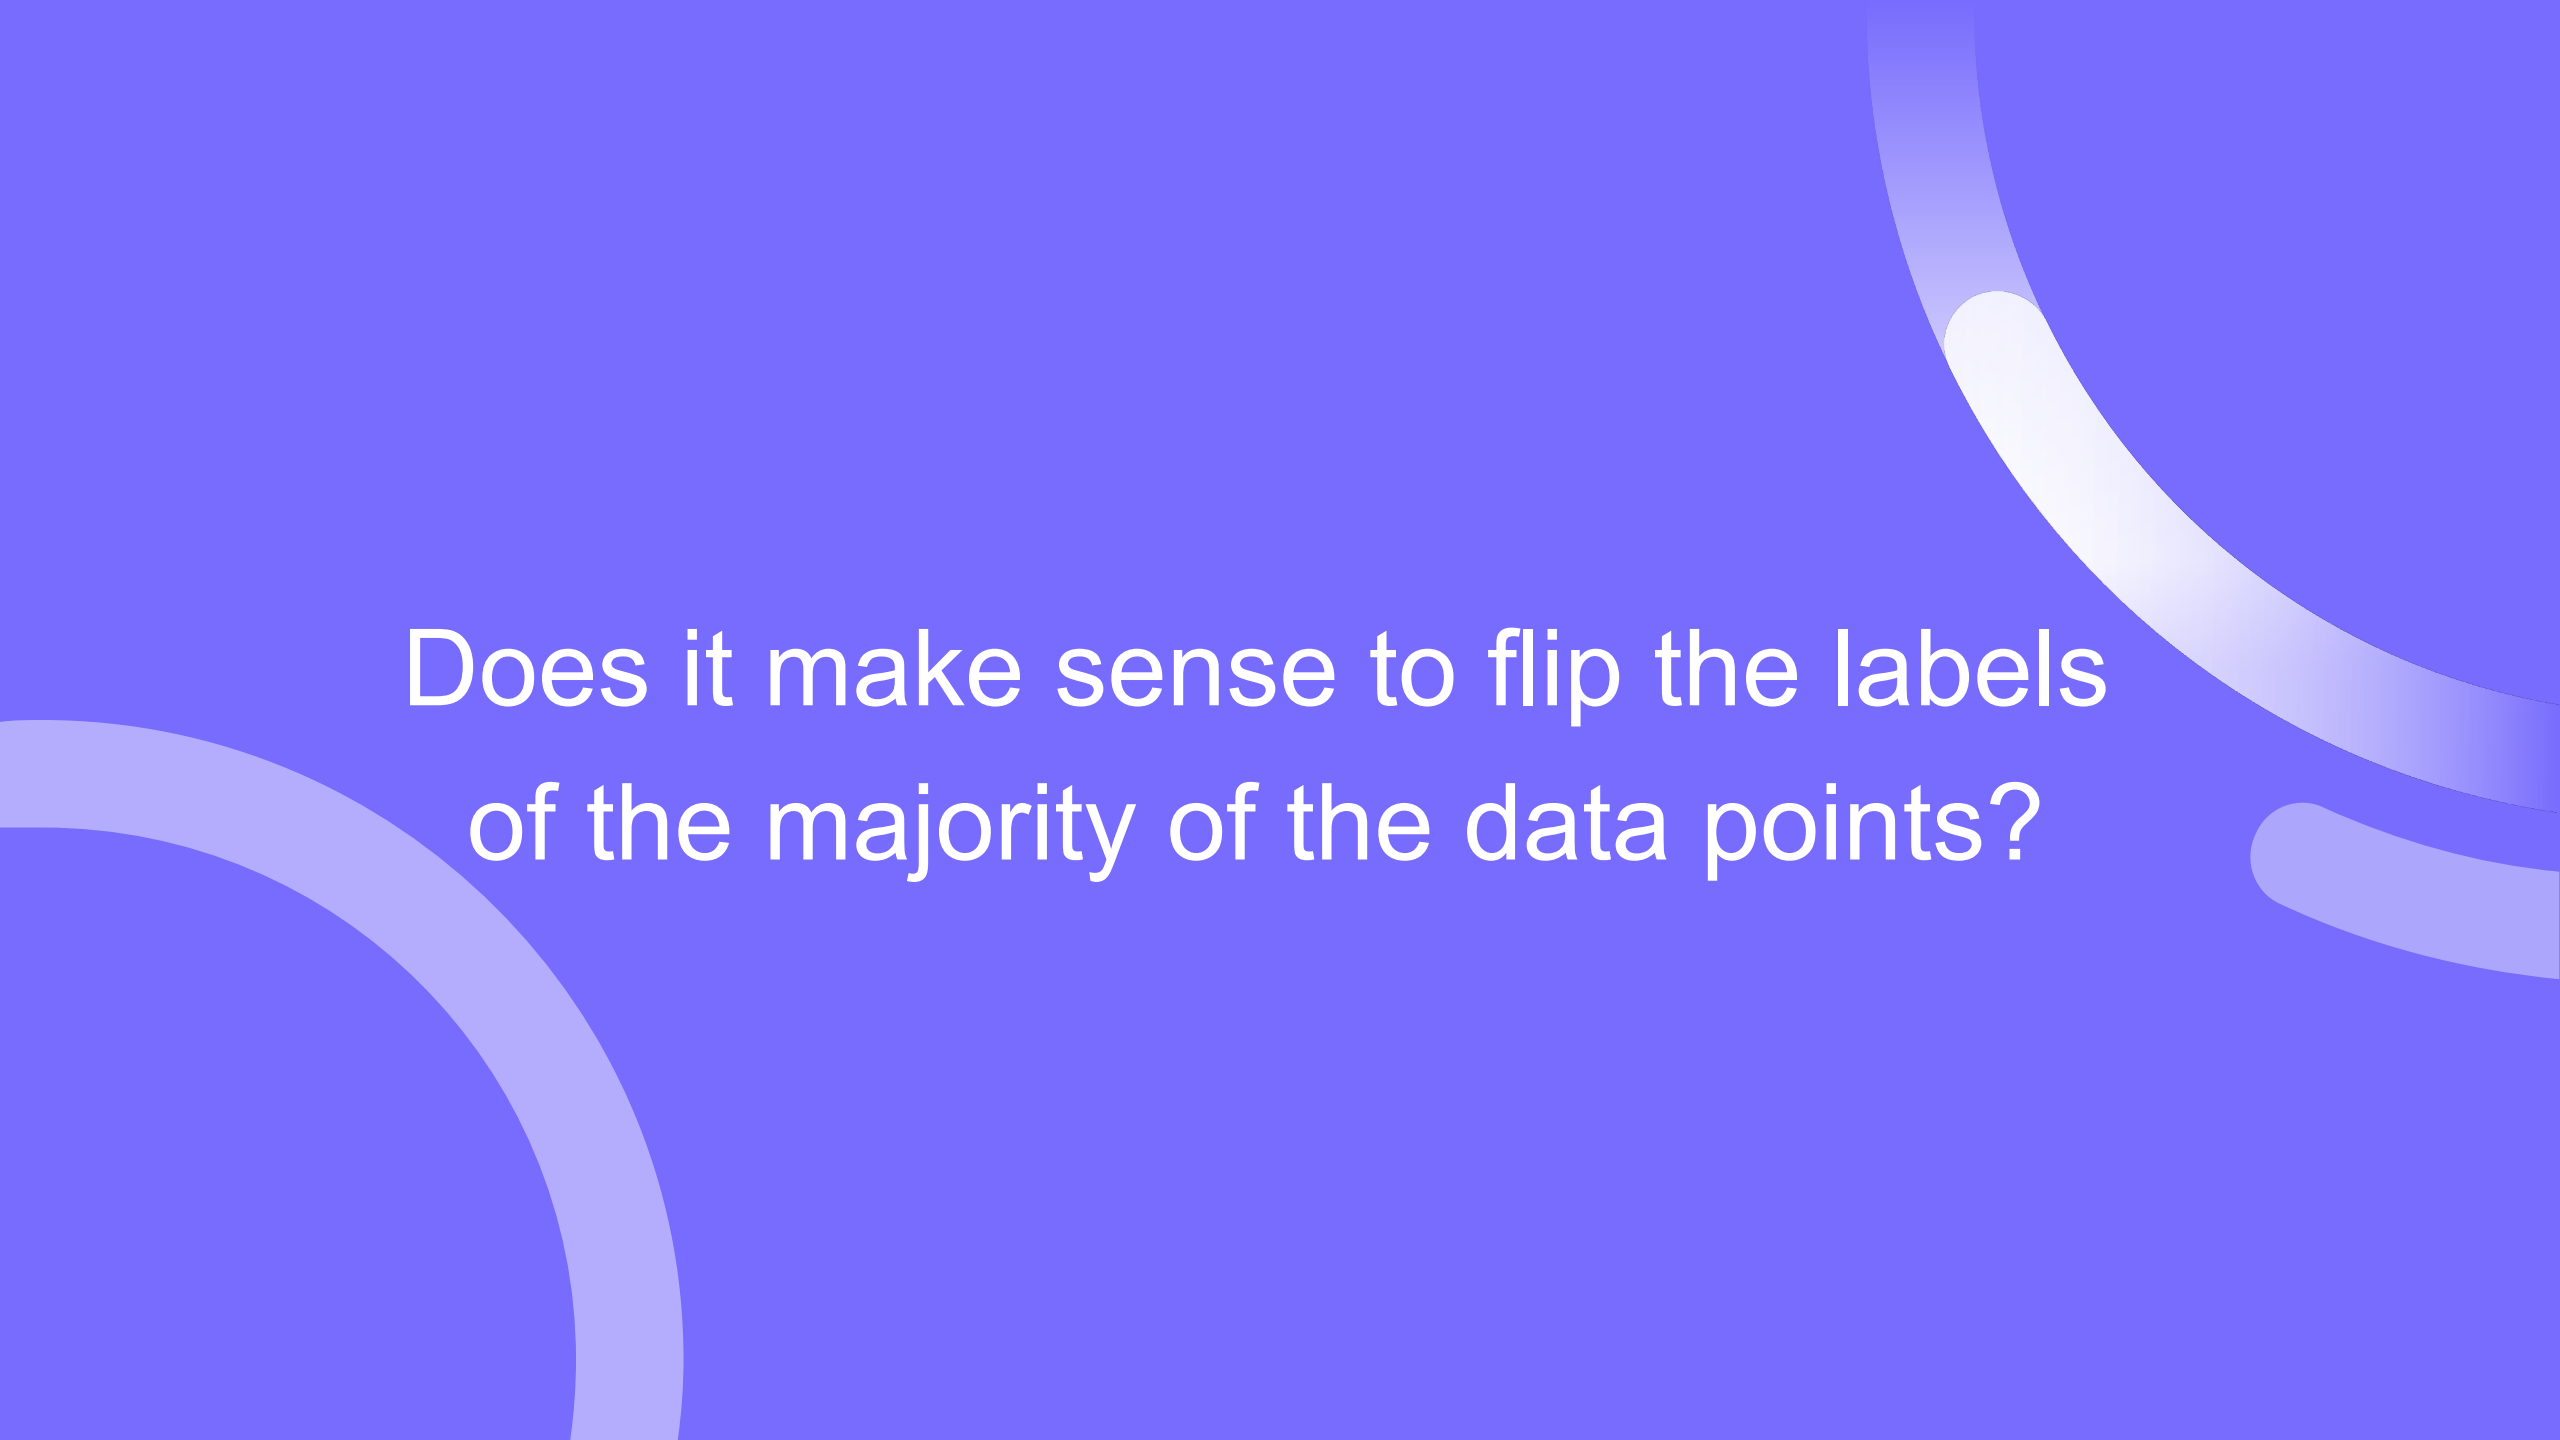


New Distributions of Target and Categorical Features



New Distributions of Continuous Features





Does it make sense to flip the labels
of the majority of the data points?